Regime Structure of Adaptive Memory Systems

Gabriel Aubin-Moreau*

March 2026

Abstract

Papers 1–6 of this series established four constraints governing adaptive memory systems and proposed a conceptual regime map with two axes: turnover rate ν and propagation strength κ. This paper maps that structure empirically in a Viability-Constrained Coupled Map Lattice (VCML).

We report seven experiments. *Primary* (1–4): (1) Sweeping ν reveals that `sg4` is *non-monotone*, peaking at intermediate turnover ($\nu \approx 0.001$–0.002) and collapsing at both extremes. (2) Sweeping κ reveals a diffusion–coherence trade-off: sg4 decreases but nonadj/adj increases with κ. (3) A 5×5 sweep maps the empirical regime structure: the adaptive regime occupies a connected interior band bounded by frozen (low ν), turnover-dominated (high ν), and fragmented (low κ) boundaries. (4) Time-series fingerprints show a *dip-and-recovery* signature unique to the adaptive regime, robust to a $2\times$ increase in perturbation magnitude. *Validation* (5–7): (5) Logistic regression zone prediction peaks in the adaptive regime, confirming regime boundaries with an orthogonal metric. (6) Ablating the copy-forward loop (SEED_BETA = 0) reduces the adaptive peak by up to 52%, causally validating the proposed mechanism. (7) Scaling from $N = 200$ to $N = 800$ sites leaves the regime structure qualitatively unchanged, ruling out finite-size artifacts.

These results ground the constraint framework of Papers 1–6 in directly measurable dynamical structure.

1 Introduction

Papers 0–6 of this series progressively constructed a constraint-based framework for adaptive memory. Paper 0 demonstrated spontaneous spatial self-organization. Papers 1–4 derived four constraints—capacity, bandwidth, propagation, and stability. Paper 5 unified them into a feasibility condition. Paper 6 interpreted the framework in the context of AI architectures, proposing a conceptual phase diagram organized by two axes: turnover rate ν and propagation strength κ.

That diagram remained theoretical. This paper measures it.

The core question. What dynamical regimes exist in an adaptive memory substrate when ν and κ are varied? Is the adaptive regime a point or an extended region? Are regime transitions sharp or gradual? How do structural and adaptability metrics co-vary across the plane?

*Code and papers: `https://github.com/AccidentalGenius101/adaptive-memory-theory` DOI: 10.5281/zenodo.XXXXXXX

Experimental substrate. We use a VCML with 400 active sites, `HS` = 2, and four spatial zones receiving structured wave input. Two control parameters:

- ν = `DEATH_P`: Bernoulli collapse probability per site per step, controlling turnover rate;
- κ = `KAPPA`: field diffusion rate, controlling how strongly each site's `fieldM` is blended with its neighbors' per step.

Two primary metrics:

- `sg4`: mean pairwise L2 distance between zone-mean `fieldM` (zone differentiation);
- `nonadj/adj`: ratio of non-adjacent to adjacent zone distances, a scale-invariant test for spatial location encoding.

All experiments use five seeds and 2000 steps with a pattern shift at step 1000.

2 Background

2.1 The two axes

Turnover ν. In VCSM, each collapse event rebirths a cell from local debris, seeding it with `fieldM` from a neighboring donor. This copy-forward loop is the mechanism by which spatial structure propagates intergenerationally. The rate ν determines:

- At $\nu \to 0$ (frozen): cells never rebirth. `fieldM` accumulates via direct consolidation only, without copy-forward amplification. Structure is rigid: strongly encoded Phase 1 patterns resist overwriting.
- At $\nu \to \infty$ (turnover-dominated): cells rebirth faster than the consolidation gate can fire. `fieldM` is continuously erased before stabilizing.
- At intermediate ν (adaptive): enough rebirths per wave cycle for copy-forward to amplify new input across zones, but slow enough for the stability gate (`STABLE_STEPS` = 10) to write.

Propagation κ. Field diffusion blends each site's `fieldM` toward its neighbors:

$$\Delta\texttt{fieldM}[i] = \kappa \cdot \big(\langle \texttt{fieldM}[\text{nbs}]\rangle - \texttt{fieldM}[i]\big) \tag{1}$$

At $\kappa = 0$: each site's field evolves independently. Zone-level coherence can form only from the shared input signal, not from spatial smoothing. This is the fragmented regime. At moderate κ: diffusion propagates local updates across zone boundaries, building coherent spatial gradients.

2.2 Predicted regime map

Regime	ν	κ
Frozen	low	any
Turnover-dom	high	any
Fragmented	mid	low
Adaptive	mid	mid–high

3 Experiment 1: Turnover Sweep

3.1 Protocol

$\kappa = 0.020$, `SEED_BETA` $= 0.25$ fixed. Sweep $\nu \in \{0.0001, 0.0003, 0.001, 0.002, 0.005, 0.010, 0.020, 0.040, 0.080\}$. 5 seeds, 2000 steps.

3.2 Results

Table 1: Turnover sweep: `sg4` is non-monotone in ν, peaking at intermediate turnover.

ν	coll/site	sg4	Regime
0.0001	0.00018	223	frozen
0.0003	0.00056	273	frozen
0.0010	0.00191	**376**	adaptive
0.0020	0.00388	**328**	adaptive
0.0050	0.00955	80	adaptive
0.0100	0.01905	24	transitional
0.0200	0.03830	4.5	turnover-dom
0.0400	0.07556	0.9	turnover-dom
0.0800	0.14865	0.2	turnover-dom

Key finding. `sg4` is non-monotone in ν. It peaks at $\nu \approx 0.001$–0.002 (coll/site $\approx$ 0.002–0.004), then collapses toward zero at high ν. This identifies the *adaptive window* as the region where the copy-forward loop fires frequently enough to amplify zone structure without erasing it.

Frozen vs. adaptive. Frozen ($\nu \leq 0.0003$) shows moderate `sg4` (223–273), lower than the adaptive peak because without copy-forward events, zone differentiation builds slowly via direct consolidation only. The frozen regime exhibits *structural rigidity*: strong resistance to pattern shifts (Experiment 4). The adaptive regime ($\nu = 0.001$–0.002) achieves the highest `sg4` precisely because copy-forward amplifies new consolidation events at each rebirth.

Remark 1 (Optimal turnover). The empirical peak at coll/site ≈ 0.002–0.004 is consistent with prior findings (Papers V87, V88, V93 in the experimental log): too little turnover starves the copy-forward loop; too much erases field faster than the consolidation gate fires.

4 Experiment 2: Propagation Sweep

4.1 Protocol

$\nu = 0.005$, `SEED_BETA` $= 0.25$ fixed. Sweep $\kappa \in \{0.000, 0.002, 0.005, 0.010, 0.020, 0.040, 0.080, 0.150\}$. 5 seeds, 2000 steps.

4.2 Results

Key finding. The two metrics diverge: `sg4` *decreases* with κ (stronger diffusion homogenizes zone-mean magnitudes) while `nonadj/adj` *increases* (stronger diffusion stabilizes coherent

Table 2: Propagation sweep: nonadj/adj increases with κ even as sg4 decreases, revealing a diffusion-coherence trade-off.

κ	sg4	nonadj/adj	Regime
0.000	50.3	1.76	fragmented
0.002	45.6	1.70	weakly coherent
0.005	42.5	1.69	weakly coherent
0.010	39.4	1.82	coherent
0.020	35.7	2.01	coherent
0.040	32.5	2.09	coherent
0.080	28.2	2.10	coherent
0.150	22.6	2.12	coherent

spatial gradients by reducing local variance). This reveals a structural trade-off:

- Raw zone separation (`sg4`) reflects field magnitude, which is diluted by diffusion.
- Location encoding quality (`nonadj/adj`) reflects gradient shape, which is stabilized and made more coherent by diffusion.

At $\kappa = 0$, nonadj/adj $= 1.76 > 1$: even without diffusion, the shared wave input creates some zone structure. But the ratio remains lower than in the coherent regime (2.09–2.12), confirming that diffusion strengthens spatial encoding.

Remark 2 (Scale-free metric). `nonadj/adj` > 1 is a scale-free test for genuine location encoding: if zones encode position, non-adjacent zones must differ more than adjacent ones regardless of absolute magnitude. Diffusion increases this ratio by building smoother zone-level gradients that are inherently more spatial in character.

5 Experiment 3: 2D Phase Map

5.1 Protocol

Sweep (ν, κ) on a 5×5 grid:

$$\nu \in \{0.0003, 0.001, 0.005, 0.015, 0.050\}$$
$$\kappa \in \{0.000, 0.005, 0.020, 0.060, 0.150\}$$

5 seeds per condition (125 total runs). Metrics: `sg4` and early adaptation score (cosine similarity at step $1200 = \texttt{SHIFT} + 200$).

5.2 Results

The map reveals three structural features:

1. **Turnover-dominated boundary is sharp.** `sg4` drops from 8–11 ($\nu = 0.015$) to 0.4–0.8 ($\nu = 0.050$) across all κ values. Above this boundary, turnover erases field faster than consolidation can write it, regardless of diffusion.

Table 3: 2D phase map: `sg4` across the $\nu \times \kappa$ plane. Regime boundaries visible as row-wise (turnover-dominated boundary) and column-wise (diffusion effect) transitions. Adaptive window at $\nu \in [0.001, 0.005]$.

$\nu \setminus \kappa$	0.000	0.005	0.020	0.060	0.150
0.0003	279	271	273	273	276
0.0010	397	383	**376**	375	361
0.0050	99	88	80	68	49
0.0150	11	9.5	8.2	6.5	4.7
0.0500	0.8	0.7	0.6	0.5	0.4

2. **Adaptive peak is non-monotone in ν.** The highest `sg4` values (376–397) occur at $\nu = 0.001$, not at the lowest $\nu = 0.0003$ (273–279). The copy-forward loop, which fires on each rebirth, amplifies zone structure during Phase 1 encoding. Frozen systems build structure more slowly.

3. **Diffusion reduces sg4 but not coherence.** Across all ν rows, `sg4` decreases with κ: from 279 to 276 at $\nu = 0.0003$, from 397 to 361 at $\nu = 0.001$, from 99 to 49 at $\nu = 0.005$. Consistent with Experiment 2, this reflects magnitude homogenization, not degradation of spatial encoding.

Remark 3 (Feasible region). The feasible region—where `sg4` is above the noise floor (`sg4` > 10) and adaptation is possible—corresponds to $\nu \in [0.001, 0.015]$ across all κ values tested. This is a connected interior band in the 2D plane, bounded above and below by the regime transitions.

6 Experiment 4: Regime Fingerprints

6.1 Protocol

One representative seed (seed 42) per regime, 2000 steps, snapshots every 50 steps. Parameters:

Regime	ν	κ	SEED_BETA	WAVE_AMP
Frozen	0.0001	0.020	0.25	0.3
Adaptive	0.0050	0.020	0.25	0.3
Turnover-dom	0.0600	0.020	0.25	0.3
Fragmented	0.0050	0.000	0.00	0.3
Adaptive-2×	0.0050	0.020	0.25	0.6

6.2 Results

Key measurements at four timepoints ($t = 500$, 1000, 1200, 2000):

Each regime exhibits a qualitatively distinct signature:

Frozen. `sg4` rises monotonically during Phase 1 (up to 1668 at $t = 1000$) and remains high after the shift. The adapt_B score is near zero at $t = 1200$ (+0.04), indicating that the system has not yet encoded the new Phase 2 pattern despite 200 steps of Phase 2 perturbation.

Table 4: Regime fingerprints: qualitative time-series signatures. Pattern shifts at $t = 1000$. adapt_B is cosine similarity of fieldM contrast with Phase 2 direction at the given timestep. Adaptive-2× uses doubled wave amplitude (0.6 vs. 0.3); all other parameters identical to Adaptive.

Regime	t	sg4	coll/site	nonadj/adj	adapt_B
Frozen	500	685	0.0001	2.84	+0.00
	1000	1668	0.0002	2.80	+0.00
	1200	1526	0.0003	2.70	+0.04
	2000	563	0.0001	2.40	+0.99
Adaptive	500	55	0.0091	2.69	+0.00
	1000	55	0.0090	2.52	+0.00
	1200	8.3	0.0079	1.20	**+0.62**
	2000	55	0.0079	2.22	+0.99
Turnover-dom	500	0.63	0.118	2.82	−0.01
	1000	0.32	0.117	2.40	−0.01
	1200	1.33	0.118	2.24	+0.99
	2000	0.65	0.116	1.65	+0.99
Fragmented	500	60	0.0091	2.71	+0.00
	1000	56	0.0090	1.72	+0.00
	1200	15	0.0079	1.05	+0.74
	2000	58	0.0079	1.84	+0.99
Adaptive-2×	500	109	0.0091	2.69	+0.00
	1000	110	0.0090	2.52	+0.00
	1200	16.5	0.0079	1.20	**+0.62**
	2000	110	0.0079	2.22	+0.99

This slow adaptation is the hallmark of the frozen regime: without copy-forward rebirths to propagate new field values rapidly, the system must update zone by zone through direct consolidation alone.

Adaptive. `sg4` is moderate but stable during Phase 1 (≈ 55). At the pattern shift, `sg4` drops sharply (from 55 to 8.3 by $t = 1200$) as the new pattern overwrites old zone structure. Critically, adapt_B $= +0.62$ at $t = 1200$—already well-aligned with the new pattern after only 200 steps. By $t = 2000$, `sg4` has recovered (≈ 55) and adapt_B $= +0.99$: full adaptation. The *dip-and-recovery* in `sg4` at the shift is the unique signature of the adaptive regime.

Turnover-dominated. `sg4` ≈ 0 throughout. The system never builds stable zone structure. High coll/site (≈ 0.12) confirms continuous field erasure. The adapt_B $= +0.99$ at $t = 1200$ reflects not structured encoding but rather the immediate local response to perturbation: each site's field is wiped and immediately re-seeded by the current wave, with no zone-level memory.[1]

[1] We use the term *turnover-dominated* rather than *chaotic*, which in dynamical systems refers to deterministic sensitive dependence on initial conditions. The high-ν regime is better described as a memory-erasure regime: it is the consolidation gate that fails, not determinism.

Fragmented. With $\kappa = 0$ and SEED_BETA $= 0$, no spatial propagation mechanism operates. `sg4` is moderate (similar to adaptive) but nonadj/adj drops to 1.05 at $t = 1200$—near 1.0, indicating that the structure that does form is not spatial: adjacent and non-adjacent zones differ by the same amount. The system relearns locally but cannot maintain zone-level gradients without propagation.

Adaptive-2×. Doubling the wave amplitude (0.6 instead of 0.3) scales the absolute `sg4` values (110 vs. 55 during Phase 1) but leaves the qualitative fingerprint unchanged: dip from 110 to 16.5 at $t = 1200$, adapt_B $= +0.62$ (identical to the standard adaptive), recovery to 110 by $t = 2000$. This confirms that the dip-and-recovery signature is a property of the regime, not an artifact of a specific perturbation magnitude.

Remark 4 (Dip-and-recovery as adaptive test). The signature `sg4` dip at the pattern shift is a discriminating test for the adaptive regime: it indicates that the system is actively *rewriting* zone structure (old pattern erased, new pattern encoded) rather than rigidly holding the old pattern (frozen) or having no structure to erase (turnover-dominated, fragmented). The robustness test (Adaptive-2×) shows the signature persists across a 2× change in perturbation magnitude.

7 Robustness and Validation

Three additional experiments address reviewer-level concerns about metric dependence, mechanism attribution, and finite-size artifacts.

7.1 Classifier Accuracy (Experiment 5)

Motivation. All primary results rely on `sg4`, a single distance-based metric. Could the regime boundaries be artifacts of that particular choice?

Protocol. We train a multinomial logistic regression classifier to predict zone membership (0–3) from the `fieldM` vector of each site at the end of the run. 4-class, 400 balanced samples, 5-fold cross-validation, 5 seeds. Chance level: 0.25 (100 sites per zone, 4 balanced classes).

Table 5: Classifier accuracy: zone prediction from `fieldM` (5 seeds, 5-fold CV). Chance level: 0.25.

ν	coll/site	Accuracy	Regime
0.0001	0.00018	0.424	frozen
0.0003	0.00056	0.349	frozen
0.0010	0.00191	0.394	adaptive
0.0020	0.00388	**0.572**	adaptive
0.0050	0.00955	**0.598**	adaptive
0.0100	0.01905	0.567	transitional
0.0200	0.03830	0.430	turnover-dom
0.0400	0.07556	0.501	turnover-dom
0.0800	0.14865	0.461	turnover-dom

Result. Accuracy exceeds chance at all ν values, reflecting that zone structure is partially encoded even in non-adaptive regimes (the wave input itself creates transient zone correlations

that persist in the final fieldM). This is an important distinction: classifier accuracy reflects instantaneous signal correlations present in the final snapshot, whereas `sg4` measures persistent spatial memory that survives across perturbation phases. Critically, accuracy peaks in the adaptive regime (ν = 0.002–0.005: 0.57–0.60), consistent with `sg4` identifying the same window. The regime pattern is independent of the choice of metric.

7.2 Copy-Forward Ablation (Experiment 6)

Motivation. We attribute the adaptive peak to the *copy-forward loop*: at each rebirth, `fieldM` from a neighboring donor is seeded into the new cell. Could the peak arise from other mechanisms, such as stochastic mixing?

Protocol. We ablate the copy-forward loop by setting SEED_BETA = 0: births use pure debris noise, receiving no field inheritance. Field diffusion (KAPPA = 0.020) remains active, isolating intergenerational from within-lifetime propagation. We run the same ν sweep as Experiment 1 under both conditions (5 seeds each).

Table 6: Copy-forward ablation: `sg4` with (standard, SEED_BETA=0.25) and without (ablated, SEED_BETA=0) field seeding at birth. KAPPA=0.020 active in both conditions.

ν	sg4 std	sg4 abl	Ratio	Regime
0.0001	223	210	1.06	frozen
0.0003	273	289	0.94	frozen
0.0010	**376**	**311**	1.21	adaptive
0.0020	**328**	**215**	**1.52**	adaptive
0.0050	79.8	60.2	1.33	adaptive
0.0100	24.1	18.0	1.34	transitional
0.0200	4.5	4.5	1.00	turnover-dom
0.0400	0.95	0.92	1.03	turnover-dom
0.0800	0.22	0.22	1.00	turnover-dom

Result. Copy-forward amplifies the adaptive peak by up to 52% (ratio = 1.52 at $\nu = 0.002$). In the frozen regime, copy-forward has no effect (too few rebirths); in the turnover-dominated regime, copy-forward has no effect (structure is erased before it can consolidate). The amplification is uniquely concentrated in the adaptive window, where rebirth rate is sufficient for copy-forward to contribute and slow enough for consolidation to write. This confirms copy-forward as a mechanism that amplifies adaptive structure formation, not merely a correlated feature. (The peak persists at reduced amplitude after ablation, so copy-forward amplifies rather than strictly enables the regime.)

7.3 Scaling Robustness (Experiment 7)

Motivation. All experiments use $N = 400$ active sites. Could the adaptive window shift or collapse at other system sizes?

Protocol. We vary grid height ($H = 10, 20, 40$) while keeping the column structure fixed (HALF = 20, ZONE_W = 5, 4 zones), giving $N = 200, 400, 800$ sites. Wave rate WR = 4.8 is

held constant across all sizes (same per-site wave intensity: each wave always covers 50% of sites regardless of H). Same ν sweep, 5 seeds.

Table 7: Scaling robustness: `sg4` and nonadj/adj across $N = 200, 400, 800$ at key ν values. WR = 4.8 held constant. The adaptive peak at $\nu = 0.001$ and the turnover-dominated collapse at $\nu = 0.050$ are preserved at all three sizes.

N	ν	sg4	nonadj/adj	Regime
	0.001	**424**	2.11	adaptive
200	0.005	78.0	1.70	adaptive
	0.050	0.27	1.81	turnover-dom
	0.001	**376**	2.16	adaptive
400	0.005	79.8	2.07	adaptive
	0.050	0.22	1.76	turnover-dom
	0.001	**484**	2.54	adaptive
800	0.005	101	2.49	adaptive
	0.050	0.26	1.73	turnover-dom

Result. The adaptive peak at $\nu = 0.001$ (sg4 = 424, 376, 484 for $N = 200, 400, 800$) and the turnover-dominated collapse at $\nu = 0.050$ (sg4 $\approx$ 0.22–0.27) are present at all three sizes. Nonadj/adj > 1 in the adaptive regime confirms spatial encoding. The regime structure is not a finite-size artifact.

8 Discussion

8.1 The empirical regime map

The four experiments collectively map the conceptual regime structure of Paper 6. The adaptive regime is confirmed as a connected interior region: $\nu \in [0.001, 0.015]$ with moderate-to-high κ. It is bounded by:

- The frozen boundary at low ν: detected as the onset of the sg4 peak (Experiment 1) and structural rigidity (Experiment 4);
- The turnover-dominated boundary at high $\nu \gtrsim 0.020$: detected as the collapse of `sg4` (Experiment 1, Table 3);
- The fragmented boundary at $\kappa \to 0$: detected as `nonadj/adj` $\to 1$ (Experiments 2, 4).

8.2 Regime violations and the four constraints

Each regime maps to a specific constraint violation:

- Frozen: violates the bandwidth constraint (Paper 2). $\nu \to 0$ means new information cannot replace old representations at a sufficient rate.

- Turnover-dominated: violates the stability constraint (Paper 4). High turnover prevents the consolidation gate from firing, so `fieldM` never stabilizes.
- Fragmented: violates the propagation constraint (Paper 3). $\kappa = 0$ eliminates spatial smoothing; the correlation length is zero.
- Adaptive: satisfies all four constraints simultaneously.

8.3 The dip-and-recovery as a diagnostic

The `sg4` dip at the pattern shift is not a failure; it is the sign that learning is occurring. A system that shows no dip at the shift has not adapted. A system that dips and does not recover has entered a turnover-dominated or fragmented regime. Only the adaptive regime shows dip-and-recovery, and only the adaptive regime recovers to match the new pattern with high nonadj/adj.

This signature can serve as a diagnostic for adaptive memory in any substrate: expose the system to a distribution shift and observe whether `sg4` transiently decreases and then recovers to a comparable level.

8.4 Robustness summary

Three additional experiments address potential objections:

- **Metric dependence** (Experiment 5): Logistic regression accuracy peaks in the same adaptive window as `sg4`, confirming the regime structure is visible to a completely different metric.
- **Mechanism attribution** (Experiment 6): Ablating the copy-forward loop (SEED_BETA = 0) reduces the adaptive peak by up to 52%, while having no effect in the frozen or turnover-dominated regimes. Copy-forward is a necessary amplifier of the adaptive peak.
- **Finite-size** (Experiment 7): The adaptive window structure persists at $N = 200$ and $N = 800$, ruling out finite-size artifact.

8.5 Limitations

Regime boundaries reported here are specific to this substrate (VCML, 400 active sites, HS=2) and one input protocol (structured zone waves). Precise boundary locations shift with system size and other hyperparameters; the structural form of the map (four regimes, interior adaptive window) is the robust prediction, not the exact ν and κ thresholds. The current map samples a 5×5 parameter grid; finer sampling could in principle reveal internal structure within the adaptive window and will be explored in future work. Testing under different wave amplitudes, frequencies, or spatial patterns also remains for future work; the shift-magnitude robustness (Experiment 4, Adaptive-2×) provides partial confidence that the signatures are input-independent. The frozen regime in VCML is not completely rigid: cells still consolidate through the normal gate. Full rigidity requires explicitly blocking consolidation, as in Paper 6 Experiment 3. The experiments here use the more natural operational definition: frozen = very low death probability, not zero plasticity.

8.6 Interpretation of Regimes

The regime terminology used in this paper should be understood as *descriptive* rather than as a claim of sharp phase transitions in the strict statistical-physics sense. The experiments reveal qualitatively distinct dynamical behaviors across the (ν, κ) plane—frozen persistence, turnover-dominated erasure, fragmented local learning, and adaptive rewrite dynamics—but the boundaries between these behaviors appear continuous rather than discontinuous at the parameter resolutions tested here. The purpose of the empirical regime map is therefore practical: to summarize the dominant dynamical patterns observed across parameter space and to guide further investigation of adaptive memory substrates. Higher-resolution parameter sweeps, alternative input statistics, and additional metrics will be useful for refining this map in future work.

9 Conclusion

We empirically mapped the dynamical regime structure of a VCML adaptive memory substrate across the turnover–propagation plane. Four regimes emerge: frozen (low ν), turnover-dominated (high ν), fragmented (low κ), and adaptive (intermediate ν, moderate κ). The adaptive regime is identified by three converging signatures: (1) a non-monotone peak in `sg4` as a function of ν; (2) a `nonadj/adj` > 1 location encoding signature; and (3) a dip-and-recovery trajectory in `sg4` at a mid-run pattern shift. The four regimes correspond to qualitatively distinct dynamical behaviors—not merely different performance levels on a continuous landscape. Experiment 4 documents the divergent fingerprint trajectories (frozen lock-in, dip-and-recovery, turnover-dominated structural collapse) that separate them regardless of metric choice. Each regime boundary corresponds to violation of one of the four fundamental constraints derived in Papers 1–5. Three robustness experiments validate the claims against reviewer concerns: an orthogonal logistic classifier confirms the regime map with a different metric; a copy-forward ablation validates copy-forward as the amplification mechanism; and scaling to $N = 800$ rules out finite-size artifacts. The empirical regime map transforms the theoretical framework into a measurable landscape.

A Reproducibility

All experiments use only `numpy`. No ML framework required.

Script	Description	Runtime
`paper7_exp1_turnover_sweep.py`	Turnover sweep (ν)	∼3 min
`paper7_exp2_propagation_sweep.py`	Propagation sweep (κ)	∼3 min
`paper7_exp3_phase_map.py`	2D phase map (125 runs)	∼10 min
`paper7_exp4_regime_fingerprints.py`	Time-series fingerprints	∼3 min
`paper7_exp5_classifier.py`	Classifier accuracy (sklearn)	∼5 min
`paper7_exp6_ablation.py`	Copy-forward ablation	∼6 min
`paper7_exp7_scaling.py`	Scaling N=200/400/800	∼8 min

Listing 1: Run all experiments (results cached after first run). Requires: numpy for Exp1-4,7; numpy + scikit-learn for Exp5.

```
cd papers/paper7_regime_structure
py paper7_exp1_turnover_sweep.py
py paper7_exp2_propagation_sweep.py
py paper7_exp3_phase_map.py
py paper7_exp4_regime_fingerprints.py
py paper7_exp5_classifier.py
py paper7_exp6_ablation.py
py paper7_exp7_scaling.py
```

Listing 2: Core VCSM step with variable nu and kappa

```
# [Calm] baseline tracker
base += BASE_BETA * (h - base)
# [Perturb] contrastive accumulation during wave
m[pert] += ALPHA_MID * (h - base)[pert];  m *= MID_DECAY
# [Consolidate] gated by calm streak
streak[pert] = 0;  streak[~pert] += 1
F[ok := streak >= SS] += FIELD_ALPHA * (m[ok] - F[ok])
F *= FIELD_DECAY
# [Diffuse] spatial smoothing with rate kappa
dF[i] = kappa * (mean(F[neighbors]) - F[i])
# [Birth] rebirth from debris + field
h_new = (1-seed_beta)*h_debris + seed_beta*F[donor]
```